

INTRODUCTION

This dataset comprises 20,000 fully synthetic customer financial profiles from India, integrating demographic attributes, income and credit indicators, and transaction-level records across diverse merchants and geographies. Generated using controlled algorithms, the data mirrors realistic financial and market behavior while maintaining complete anonymity, making it well-suited for analytical exploration, modeling, and research without privacy concerns.

IMPORTING LIBRARIES AND DATA SET LOADING

In [26]:

```
import warnings
warnings.filterwarnings("ignore")

import pandas as pd
import numpy as np

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import mean_squared_error

from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.cluster import KMeans

from sklearn.metrics import (
    r2_score,
    mean_squared_error,
    accuracy_score,
    classification_report,
    confusion_matrix,
    roc_auc_score,
    roc_curve
)
import matplotlib.pyplot as plt
import seaborn as sns
```

In [2]:

```
df = pd.read_csv(r"C:\Users\papus\OneDrive\Documents\PORTFOLIO PROJECT 2\Customer_financial_profiles.csv")
df
```

Out[2]:

	id	current_age	birth_year	birth_month	gender	address	per_capita_income	yearly_income	total_debt	credit_score	...	transaction_id	date	client_id	card_ic
0	1	64	1961	3	Female	Street 286, Indore, Madhya Pradesh - 452645	86125	1033501	280435	715	...	1251	28- 03- 2024	CLTIND00104	CRD001042
1	2	39	1986	12	Female	Street 268, Kolkata, West Bengal - 700555	68414	820965	0	834	...	9976	20- 09- 2025	CLTIND00213	CRD002132
2	3	50	1975	7	Female	Street 262, Lucknow, Uttar Pradesh - 226797	22595	271138	0	609	...	6084	02- 10- 2024	CLTIND03292	CRD032922
3	4	27	1998	2	Male	Street 265, Nagpur, Maharashtra - 440623	24158	289898	0	782	...	3350	02- 09- 2023	CLTIND05205	CRD052053
4	5	45	1980	7	Other	Street 157, Delhi, Delhi - 110965	45475	545697	117821	608	...	11725	25- 05- 2025	CLTIND04910	CRD049101
...
19995	19996	39	1986	10	Other	Street 112, Chennai, Tamil Nadu - 600341	33596	403151	142167	582	...	6718	24- 02- 2024	CLTIND04778	CRD047781
19996	19997	44	1981	6	Female	Street 197, Ahmedabad, Gujarat - 380460	66484	797812	0	647	...	15264	21- 06- 2024	CLTIND00150	CRD001502
19997	19998	43	1982	1	Female	Street 7, Indore, Madhya Pradesh - 452341	41053	492634	0	853	...	8536	05- 03- 2025	CLTIND02792	CRD027923
19998	19999	49	1976	9	Female	Street 212, Nagpur, Maharashtra - 440770	51856	622267	266802	834	...	14325	06- 11- 2024	CLTIND05714	CRD057142
19999	20000	37	1988	4	Male	Street 223, Indore, Madhya	43109	517305	130597	753	...	9705	27- 06- 2024	CLTIND00538	CRD005381

```
    id  current_age  birth_year  birth_month  gender      address  per_capita_income  yearly_income  total_debt  credit_score  ...  transaction_id  date        client_id  card_id
--  --  --  --  --  --  --  --  --  --  --  --  --  --  --  --  --  --  --  --  --  --  --
```

Pradesh -
452307

20000 rows × 21 columns

In [3]: `df.columns`

```
Out[3]: Index(['id', 'current_age', 'birth_year', 'birth_month', 'gender', 'address',
   'per_capita_income', 'yearly_income', 'total_debt', 'credit_score',
   'num_credit_cards', 'transaction_id', 'date', 'client_id', 'card_id',
   'amount', 'use_chip', 'merchant_id', 'merchant_city', 'merchant_state',
   'zip'],
  dtype='object')
```

DATA CLEANING

In [5]: `df.dtypes`

```
Out[5]: id          int64
current_age  int64
birth_year   int64
birth_month  int64
gender       object
address      object
per_capita_income  int64
yearly_income  int64
total_debt   int64
credit_score  int64
num_credit_cards  int64
transaction_id  int64
date         object
client_id    object
card_id      object
amount       float64
use_chip     object
merchant_id  object
merchant_city object
merchant_state object
zip          int64
dtype: object
```

```
In [6]: # Convert date column
df['date'] = pd.to_datetime(df['date'])

# Convert categorical columns
categorical_cols = [
    'gender',
```

```
'use_chip',
'merchant_city',
'merchant_state',
'client_id',
'card_id',
'merchant_id'
]

df[categorical_cols] = df[categorical_cols].astype('category')
```

In [7]: `df.dtypes`

```
Out[7]: id                  int64
current_age          int64
birth_year           int64
birth_month          int64
gender               category
address              object
per_capita_income    int64
yearly_income        int64
total_debt           int64
credit_score         int64
num_credit_cards    int64
transaction_id      int64
date                 datetime64[ns]
client_id            category
card_id              category
amount               float64
use_chip             category
merchant_id          category
merchant_city        category
merchant_state       category
zip                  int64
dtype: object
```

In [8]: `df.isnull().sum()`

```
Out[8]: id          0  
        current_age  0  
        birth_year   0  
        birth_month  0  
        gender       0  
        address      0  
        per_capita_income 0  
        yearly_income 0  
        total_debt    0  
        credit_score  0  
        num_credit_cards 0  
        transaction_id 0  
        date         0  
        client_id    0  
        card_id      0  
        amount        0  
        use_chip     0  
        merchant_id  0  
        merchant_city 0  
        merchant_state 0  
        zip          0  
        dtype: int64
```

```
In [10]: df.duplicated().sum()
```

```
Out[10]: np.int64(0)
```

```
In [11]: cat_cols = df.select_dtypes(include=['category', 'object']).columns  
  
for col in cat_cols:  
    print(f"\nValue counts for {col}:")  
    print(df[col].value_counts())
```

Value counts for gender:

gender
Male 9855
Female 9479
Other 666
Name: count, dtype: int64

Value counts for address:

address
Street 11, Lucknow, Uttar Pradesh - 226280 22
Street 225, Patna, Bihar - 800218 20
Street 224, Bhopal, Madhya Pradesh - 462594 20
Street 136, Delhi, Delhi - 110337 19
Street 193, Bhopal, Madhya Pradesh - 462496 19
..
Street 85, Chennai, Tamil Nadu - 600443 1
Street 36, Patna, Bihar - 800690 1
Street 284, Bhopal, Madhya Pradesh - 462823 1
Street 205, Ahmedabad, Gujarat - 380233 1
Street 148, Jaipur, Rajasthan - 302873 1
Name: count, Length: 4938, dtype: int64

Value counts for client_id:

client_id
CLTIND00062 22
CLTIND00257 20
CLTIND01752 20
CLTIND01604 19
CLTIND02908 19
..
CLTIND00077 1
CLTIND05918 1
CLTIND05929 1
CLTIND05935 1
CLTIND05945 1
Name: count, Length: 4941, dtype: int64

Value counts for card_id:

card_id
CRD029871 7
CRD054061 7
CRD049712 7
CRD003272 7
CRD044521 7
..
CRD000151 1
CRD000211 1
CRD000241 1
CRD000242 1
CRD000244 1
Name: count, Length: 10030, dtype: int64

```
Value counts for use_chip:
```

```
use_chip
```

```
Yes    13927
```

```
No     6073
```

```
Name: count, dtype: int64
```

```
Value counts for merchant_id:
```

```
merchant_id
```

```
MCH00471    63
```

```
MCH00465    62
```

```
MCH00143    60
```

```
MCH00327    60
```

```
MCH00066    57
```

```
..
```

```
MCH00165    25
```

```
MCH00499    25
```

```
MCH00171    25
```

```
MCH00118    23
```

```
MCH00418    19
```

```
Name: count, Length: 500, dtype: int64
```

```
Value counts for merchant_city:
```

```
merchant_city
```

```
Surat      1761
```

```
Bengaluru  1567
```

```
Pune       1442
```

```
Lucknow    1435
```

```
Indore     1384
```

```
Nagpur     1365
```

```
Hyderabad  1352
```

```
Mumbai     1325
```

```
Patna      1311
```

```
Jaipur     1288
```

```
Chennai    1284
```

```
Ahmedabad  1282
```

```
Bhopal     1200
```

```
Delhi      1079
```

```
Kolkata    925
```

```
Name: count, dtype: int64
```

```
Value counts for merchant_state:
```

```
merchant_state
```

```
Maharashtra 4132
```

```
Gujarat     3043
```

```
Madhya Pradesh 2584
```

```
Karnataka   1567
```

```
Uttar Pradesh 1435
```

```
Telangana    1352
```

```
Bihar        1311
```

```
Rajasthan    1288
```

```
Tamil Nadu   1284
```

```
Delhi        1079
```

```
West Bengal      925  
Name: count, dtype: int64
```

```
In [15]: ## creating Multiclass credit Tier
```

```
def credit_tier(score):  
    if score < 580:  
        return "Poor"  
    elif score < 670:  
        return "Fair"  
    elif score < 740:  
        return "Good"  
    else:  
        return "Excellent"  
  
df["Credit_Tier"] = df["credit_score"].apply(credit_tier)
```

Classification

```
In [18]: features = [  
    "yearly_income",  
    "total_debt",  
    "num_credit_cards",  
    "current_age"  
]  
  
X = df[features]  
y = df["Credit_Tier"]  
  
label_encoder = LabelEncoder()  
y_encoded = label_encoder.fit_transform(y)  
  
print("Class Mapping:")  
for i, cls in enumerate(label_encoder.classes_):  
    print(f"{cls} -> {i}")  
  
X_train, X_test, y_train, y_test = train_test_split(  
    X,  
    y_encoded,  
    test_size=0.3,  
    random_state=42,  
    stratify=y_encoded  
)  
scaler = StandardScaler()  
X_train_scaled = scaler.fit_transform(X_train)  
X_test_scaled = scaler.transform(X_test)  
  
model = LogisticRegression(  
    max_iter=1000,  
    solver="lbfgs")
```

```

)
model.fit(X_train_scaled, y_train)

y_pred = model.predict(X_test_scaled)

print("\nClassification Report:\n")
print(
    classification_report(
        y_test,
        y_pred,
        target_names=label_encoder.classes_
    )
)

print("\nConfusion Matrix:\n")
print(confusion_matrix(y_test, y_pred))

```

Class Mapping:

Excellent -> 0
 Fair -> 1
 Good -> 2
 Poor -> 3

Classification Report:

	precision	recall	f1-score	support
Excellent	0.39	0.51	0.44	1873
Fair	0.39	0.24	0.29	1788
Good	0.35	0.41	0.38	2046
Poor	1.00	0.00	0.01	293
accuracy			0.37	6000
macro avg	0.53	0.29	0.28	6000
weighted avg	0.40	0.37	0.35	6000

Confusion Matrix:

```

[[955 205 713  0]
 [597 422 769  0]
 [823 375 848  0]
 [ 82  94 116  1]]

```

INSIGHTS – Multi-Class Credit Score Classification

- 1 Model performance is realistic and leakage-free

The model achieved ~37% accuracy, which is significantly higher than the random baseline (25%) for a 4-class problem.

Performance is realistic because no target leakage was introduced and only indirect financial features were used.

2 Errors occur mainly between adjacent credit tiers

Most misclassifications happen between Excellent ↔ Good and Good ↔ Fair.

There are almost no extreme errors, such as classifying Excellent customers as Poor.

3 Fair and Good tiers are the hardest to distinguish

The Fair class shows the lowest recall and F1-score.

This indicates strong overlap in income, debt, and card usage patterns between Fair and Good customers.

4 Poor-credit customers are rarely flagged (conservative behavior)

The model predicts the Poor class conservatively, resulting in low recall for that segment.

This reduces false positives but increases the risk of missing some high-risk customers.

CREDIT SCORE MODELLING

Target: credit_score (continuous) Features: indirect financial & demographic variables

```
In [28]: features = [
    "yearly_income",
    "total_debt",
    "num_credit_cards",
    "current_age"
]

X = df[features]
y = df["credit_score"]

X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.3,
    random_state=42
)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

model = LinearRegression()
model.fit(X_train_scaled, y_train)
```

```

y_pred = model.predict(X_test_scaled)

r2 = r2_score(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)

coef_df = pd.DataFrame({
    "Feature": features,
    "Coefficient": model.coef_
})

print("\nModel Coefficients:")
print(coef_df.sort_values(by="Coefficient", ascending=False))
print("RMSE:", rmse)

```

Model Coefficients:

	Feature	Coefficient
0	yearly_income	6.329531
2	num_credit_cards	1.764953
3	current_age	-2.039505
1	total_debt	-19.816087

RMSE: 75.21645739841668

INSIGHTS

The regression model shows that total debt is the strongest negative driver of credit score, while income contributes positively but to a lesser extent. The model achieves an RMSE of approximately 75 points, indicating realistic predictive accuracy given the 300–900 credit score range and limited feature set.

LOAN DEFAULT PREDICTION

```

In [30]: df["Loan_Default"] = (
    (df["credit_score"] < 620) |
    (df["total_debt"] > df["yearly_income"] * 1.2)
).astype(int)

print("Loan Default Distribution:")
print(df["Loan_Default"].value_counts())

features = [
    "yearly_income",
    "total_debt",
    "num_credit_cards",
    "current_age"
]

X = df[features]
y = df["Loan_Default"]

```

```

X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.3,
    random_state=42,
    stratify=y
)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

model = LogisticRegression(max_iter=1000)
model.fit(X_train_scaled, y_train)

y_pred = model.predict(X_test_scaled)
y_prob = model.predict_proba(X_test_scaled)[:, 1]

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

print("ROC AUC:", roc_auc_score(y_test, y_prob))

print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred))

```

Loan Default Distribution:

Loan_Default

0 17123

1 2877

Name: count, dtype: int64

Classification Report:

	precision	recall	f1-score	support
0	0.86	1.00	0.92	5137
1	0.33	0.01	0.03	863
accuracy			0.85	6000
macro avg	0.60	0.50	0.47	6000
weighted avg	0.78	0.85	0.79	6000

ROC AUC: 0.6541278809969524

Confusion Matrix:

```

[[5113  24]
 [ 851  12]]

```

INSIGHTS

The loan default model achieves an ROC–AUC of ~0.65, indicating moderate ability to distinguish defaulters from non-defaulters. • The model is conservative, prioritizing the correct identification of non-defaulters and minimizing false rejections. • Due to class imbalance, recall for defaulters is low, highlighting the need for threshold tuning or cost-sensitive learning in real-world deployment.

FRAUD DETECTION

```
In [32]: amount_threshold = df["amount"].quantile(0.99)

df["Fraud"] = (
    (df["amount"] > amount_threshold) &
    (df["num_credit_cards"] >= 4)
).astype(int)

print("Fraud Distribution:")
print(df["Fraud"].value_counts())

features = [
    "amount",
    "num_credit_cards",
    "current_age",
    "credit_score"
]

X = df[features]
y = df["Fraud"]

X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.3,
    random_state=42,
    stratify=y
)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

model = LogisticRegression(max_iter=1000)
model.fit(X_train_scaled, y_train)

y_pred = model.predict(X_test_scaled)
y_prob = model.predict_proba(X_test_scaled)[:, 1]

print("\nClassification Report:\n")
print(classification_report(y_test, y_pred))

print("ROC AUC:", roc_auc_score(y_test, y_prob))
```

```
print("\nConfusion Matrix:\n")
print(confusion_matrix(y_test, y_pred))
```

Fraud Distribution:

Fraud

0 19930

1 70

Name: count, dtype: int64

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	5979
1	0.80	0.57	0.67	21
accuracy			1.00	6000
macro avg	0.90	0.79	0.83	6000
weighted avg	1.00	1.00	1.00	6000

ROC AUC: 0.9992752411217037

Confusion Matrix:

```
[[5976  3]
 [ 9 12]]
```

INSIGHTS

"The fraud model shows strong separation between fraudulent and non-fraudulent transactions; however, the near-perfect ROC-AUC indicates potential feature leakage due to the synthetic definition of fraud. Therefore, this model is treated as an exploratory baseline rather than a production-ready fraud system."

Customer Segmentation

```
In [33]: X = df[['
    "yearly_income",
    "total_debt",
    "num_credit_cards",
    "credit_score"
]]
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

kmeans = KMeans(n_clusters=4, random_state=42)
df["Customer_Segment"] = kmeans.fit_predict(X_scaled)

df[["Customer_Segment"]].value_counts()
```

```
Out[33]: Customer_Segment
3           6162
2           5983
1           5283
0           2572
Name: count, dtype: int64
```

```
In [34]: df.groupby("Customer_Segment")[[  
    "yearly_income",  
    "total_debt",  
    "credit_score",  
    "num_credit_cards"  
]].mean().round(2)
```

```
Out[34]:
```

Customer_Segment	yearly_income	total_debt	credit_score	num_credit_cards
0	1216281.96	479771.80	665.24	3.54
1	486886.59	59023.09	774.49	2.15
2	489681.38	99820.30	642.35	2.15
3	984420.37	88890.43	716.29	3.92

INSIGHTS

- Customer segmentation reveals four distinct financial profiles ranging from prime low-risk customers to high-income but over-leveraged individuals.
- Segment 1 represents the most creditworthy customers with high credit scores and minimal debt, suitable for premium financial products.
- Segment 0, despite high income, exhibits elevated debt levels, indicating potential over-leverage and higher credit risk.
- These segments enable differentiated strategies for risk monitoring, credit approval, and targeted marketing.

TRANSACTION CLASSIFICATION

```
In [35]: def classify_transaction(amount):  
    if amount < 2000:  
        return 0  
    elif amount < 10000:  
        return 1  
    else:  
        return 2  
  
df["Transaction_Type"] = df["amount"].apply(classify_transaction)  
  
print("Transaction Type Distribution:")  
print(df["Transaction_Type"].value_counts())  
print("\n")
```

```
features = [
    "amount",
    "yearly_income",
    "credit_score",
    "num_credit_cards"
]

X = df[features]
y = df["Transaction_Type"]

X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.30,
    random_state=42,
    stratify=y
)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

model = LogisticRegression(max_iter=1000)
model.fit(X_train_scaled, y_train)

y_pred = model.predict(X_test_scaled)

print("CLASSIFICATION REPORT:\n")
print(classification_report(y_test, y_pred))

print("CONFUSION MATRIX:\n")
print(confusion_matrix(y_test, y_pred))

coef_df = pd.DataFrame({
    "Feature": features,
    "Coefficient_Class_0": model.coef_[0],
    "Coefficient_Class_1": model.coef_[1],
    "Coefficient_Class_2": model.coef_[2]
})

print("\nMODEL COEFFICIENTS:")
print(coef_df)
```

Transaction Type Distribution:

Transaction_Type

1 12529

0 7265

2 206

Name: count, dtype: int64

CLASSIFICATION REPORT:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	2179
1	1.00	1.00	1.00	3759
2	1.00	0.90	0.95	62
accuracy			1.00	6000
macro avg	1.00	0.97	0.98	6000
weighted avg	1.00	1.00	1.00	6000

CONFUSION MATRIX:

```
[[2177  2  0]
 [ 6 3753  0]
 [ 0  6 56]]
```

MODEL COEFFICIENTS:

	Feature	Coefficient_Class_0	Coefficient_Class_1	\
0	amount	-18.434473	6.891552	
1	yearly_income	0.018008	-0.041459	
2	credit_score	-0.030510	0.006200	
3	num_credit_cards	-0.000635	0.023789	

	Coefficient_Class_2
0	11.542921
1	0.023451
2	0.024309
3	-0.023154

INSIGHTS

Transactions were categorized into low, medium, and high-value groups using a rule-based definition derived from transaction amount.

A multiclass logistic regression model was trained to validate and replicate these transaction categories, achieving near-perfect accuracy due to the deterministic relationship between transaction amount and class labels.

Model coefficients confirm that transaction amount is the dominant driver of classification, while customer attributes play a secondary role.

This approach demonstrates how rule-based transaction logic can be translated into a machine learning framework for scalable transaction monitoring.

RISK MODELLING

```
In [37]: df["High_Risk"] = (
    (df["credit_score"] < 620) |
    (df["total_debt"] > df["yearly_income"] * 1.2)
).astype(int)

print("Risk Distribution:")
print(df["High_Risk"].value_counts())
print()
features = [
    "yearly_income",
    "total_debt",
    "num_credit_cards",
    "current_age"
]

X = df[features]
y = df["High_Risk"]

X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.30,
    random_state=42,
    stratify=y
)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

risk_model = LogisticRegression(max_iter=1000)
risk_model.fit(X_train_scaled, y_train)

risk_prob = risk_model.predict_proba(X_test_scaled)[:, 1]

risk_pred = (risk_prob >= 0.30).astype(int)

print("ROC AUC:", roc_auc_score(y_test, risk_prob))
print("\nCONFUSION MATRIX:\n")
print(confusion_matrix(y_test, risk_pred))

print("\nCLASSIFICATION REPORT:\n")
print(classification_report(y_test, risk_pred))

risk_output = X_test.copy()
risk_output["Risk_Probability"] = risk_prob

risk_output["Risk_Band"] = pd.cut(
```

```

risk_output["Risk_Probability"],
bins=[0, 0.20, 0.40, 1.00],
labels=["Low Risk", "Medium Risk", "High Risk"]
)

print("\nRisk Band Distribution:")
print(risk_output["Risk_Band"].value_counts())

```

Risk Distribution:

```

High_Risk
0    17123
1     2877
Name: count, dtype: int64

```

ROC AUC: 0.6541278809969524

CONFUSION MATRIX:

```

[[5000 137]
 [ 787  76]]

```

CLASSIFICATION REPORT:

	precision	recall	f1-score	support
0	0.86	0.97	0.92	5137
1	0.36	0.09	0.14	863
accuracy			0.85	6000
macro avg	0.61	0.53	0.53	6000
weighted avg	0.79	0.85	0.80	6000

Risk Band Distribution:

```

Risk_Band
Low Risk      5098
Medium Risk    820
High Risk       82
Name: count, dtype: int64

```

INSIGHTS

The dataset contains ~15% high-risk customers, reflecting a realistic credit risk distribution.

The model achieves a ROC–AUC of ~0.65, indicating a meaningful ability to separate high-risk and low-risk customers using basic financial attributes.

Overall accuracy (~85%) is driven by class imbalance and therefore is not the primary evaluation metric for this model.

The model demonstrates high recall for low-risk customers (97%), ensuring that most creditworthy customers are correctly identified.

Recall for high-risk customers is intentionally low (9%), indicating a conservative risk strategy that minimizes false rejections.

When customers are flagged as high risk, the model shows reasonable precision (36%), meaning flagged cases are materially riskier than average.

The confusion matrix shows very few false positives, which is important in credit approval workflows where rejecting good customers is costly.


Risk probabilities are successfully translated into Low, Medium, and High risk bands, enabling actionable decision-making rather than raw predictions.

The majority of customers fall into the Low Risk band, with progressively fewer customers in Medium and High Risk bands, which aligns with real-world credit portfolios.

The model serves as a baseline, explainable credit risk framework that can be further improved through threshold tuning or additional behavioral features.

INCOME - DEBT CORRELATION

```
In [38]: corr_data = df[[  
    "yearly_income",  
    "total_debt",  
    "credit_score",  
    "num_credit_cards",  
    "current_age"  
]]  
  
corr_matrix = corr_data.corr()  
  
plt.figure(figsize=(8, 6))  
sns.heatmap(  
    corr_matrix,  
    annot=True,  
    fmt=".2f",  
    cmap="coolwarm",  
    linewidths=0.5  
)  
plt.title("Correlation Heatmap: Income, Debt & Credit Attributes")  
plt.show()
```



INSIGHTS

The heatmap shows a positive but weak-to-moderate correlation between yearly income and total debt.

Higher income does not directly imply lower financial risk, as debt levels vary significantly across income groups.

Credit score exhibits negative correlation with total debt, indicating higher debt burdens are associated with lower credit scores.

No single variable dominates the correlation structure, justifying the use of multi-feature risk models instead of rule-based decisions.

SOLVING SOME BUISNESS PROBLEM


Q.1::Which customer segments contribute the most to overall credit risk in the portfolio?

```
In [40]: # Contribution of each customer segment to overall credit risk

# Calculate risk contribution by segment
segment_risk = (
    df.groupby("Customer_Segment")["High_Risk"]
        .mean()
        .sort_values(ascending=False)
)
print(segment_risk)
# Plot
plt.figure(figsize=(8, 5))
segment_risk.plot(
    kind="bar",
    color=["red", "orange", "yellow", "green"]
)

plt.xlabel("Customer Segment")
plt.ylabel("Proportion of High-Risk Customers")
plt.title("Credit Risk Contribution by Customer Segment")
plt.xticks(rotation=0)
plt.grid(axis="y", alpha=0.3)
plt.show()
```

```
Customer_Segment
2    0.295170
0    0.285381
3    0.061181
1    0.000000
Name: High_Risk, dtype: float64
```



Customer Segment 2 and Segment 0 contribute the highest credit risk, with nearly 30% of customers in these segments classified as high risk.

Segment 1 shows negligible risk, while Segment 3 has relatively low risk, indicating clear differentiation in risk concentration across segments.

Q.2: Does higher income consistently translate into lower credit risk, or are there high-income high-risk customers?

```
In [41]: # Income vs Credit Risk (High-Risk vs Low-Risk)
```

```
plt.figure(figsize=(8, 5))

df.boxplot(
    column="yearly_income",
    by="High_Risk",
    grid=False,
    patch_artist=True,
    boxprops=dict(facecolor="skyblue"),
    medianprops=dict(color="red")
)


plt.xlabel("High Risk (0 = Low Risk, 1 = High Risk)")
```

```

plt.ylabel("Yearly Income")
plt.title("Income Distribution by Credit Risk")
plt.suptitle("")
plt.show()

```

<Figure size 800x500 with 0 Axes>



Higher income does not consistently translate into lower credit risk, as there is substantial overlap in income distributions between high-risk and low-risk customers.

High-income high-risk customers exist, indicating that debt levels and credit behavior play a more critical role than income alone in determining risk.

Q.3::What proportion of customers fall into low, medium, and high-risk bands, and how does this impact approval decisions?

```

In [47]: # Creating a Risk Band

df["Risk_Band"] = df["High_Risk"].map({
    0: "Low / Medium Risk",
    1: "High Risk"
})

# Proportion of customers in each Risk Band
risk_counts = df["Risk_Band"].value_counts(normalize=True) * 100
print(risk_counts)

```


```

# Plot
plt.figure(figsize=(8, 5))
risk_counts.plot(
    kind="bar",
    color=["green", "red"]
)

plt.xlabel("Risk Band")
plt.ylabel("Percentage of Customers (%)")
plt.title("Customer Distribution by Risk Band")
plt.xticks(rotation=0)
plt.grid(axis="y", alpha=0.3)
plt.show()

```

Risk_Band
Low / Medium Risk 85.615
High Risk 14.385
Name: proportion, dtype: float64



The majority of customers ($\approx 85.6\%$) fall into the Low/Medium Risk band, enabling high approval rates with minimal portfolio risk.

A smaller but significant segment ($\approx 14.4\%$) is classified as High Risk, indicating the need for stricter approval checks or manual review to control potential losses.

Q.4::Which financial attributes are the strongest indicators of potential loan default?

```
In [48]: features = [
    "yearly_income",
    "total_debt",
    "num_credit_cards",
    "current_age"
]

X = df[features]
y = df["High_Risk"]

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)


model = LogisticRegression(max_iter=1000)
model.fit(X_scaled, y)

coef_df = pd.DataFrame({
    "Feature": features,
    "Impact": model.coef_[0]
}).sort_values(by="Impact", ascending=False)
print(coef_df)

# Plot
plt.figure(figsize=(8, 5))
plt.barh(
    coef_df["Feature"],
    coef_df["Impact"],
    color=["red" if v > 0 else "green" for v in coef_df["Impact"]]
)

plt.xlabel("Impact on Default Risk")
plt.title("Financial Attributes Influencing Loan Default Risk")
plt.grid(axis="x", alpha=0.3)
plt.show()
```

	Feature	Impact
1	total_debt	0.513207
3	current_age	0.102020
0	yearly_income	-0.162236
2	num_credit_cards	-0.165469



Total debt is the strongest positive indicator of loan default risk, with higher debt levels significantly increasing the likelihood of default.

Yearly income and number of credit cards show negative impact, indicating that stronger income and controlled credit exposure reduce default risk, while age has a comparatively minor influence.

Q.5: How concentrated is risk across the customer base, and do a small number of customers drive most of the risk?

```
In [49]: # Risk concentration analysis: cumulative contribution of high-risk customers

# Sort customers by risk (High_Risk)
risk_sorted = df.sort_values("High_Risk", ascending=False)

# Cumulative calculations
risk_sorted["cum_customers"] = (
    np.arange(1, len(risk_sorted) + 1) / len(risk_sorted)
)


risk_sorted["cum_risk"] = (
    risk_sorted["High_Risk"].cumsum() / risk_sorted["High_Risk"].sum()
```

```
)  
print(risk_sorted)  
# Plot cumulative risk concentration curve  
plt.figure(figsize=(8, 5))  
  
plt.plot(  
    risk_sorted["cum_customers"],  
    risk_sorted["cum_risk"],  
    color="red",  
    linewidth=2,  
    label="Cumulative Risk Contribution"  
)  
  
plt.plot(  
    [0, 1],  
    [0, 1],  
    linestyle="--",  
    color="gray",  
    label="Equal Risk Distribution"  
)  
  
plt.xlabel("Cumulative % of Customers")  
plt.ylabel("Cumulative % of Risk")  
plt.title("Risk Concentration Across Customer Base")  
plt.legend()  
plt.grid(alpha=0.3)  
plt.show()
```

	id	current_age	birth_year	birth_month	gender	\	
8258	8259	31	1994	3	Male		
6665	6666	33	1992	6	Female		
8747	8748	65	1960	12	Female		
8527	8528	40	1985	7	Female		
2713	2714	36	1989	9	Female		
...		
27	28	33	1992	11	Female		
28	29	19	2006	8	Male		
29	30	51	1974	11	Female		
30	31	31	1994	1	Male		
17499	17500	65	1960	3	Male		
			address	per_capita_income	\		
8258			Street 72, Mumbai, Maharashtra - 400244	41797			
6665			Street 137, Bengaluru, Karnataka - 560677	50222			
8747			Street 189, Patna, Bihar - 800181	118676			
8527			Street 92, Nagpur, Maharashtra - 440594	42118			
2713			Street 149, Ahmedabad, Gujarat - 380797	65624			
...					
27			Street 97, Pune, Maharashtra - 411506	69372			
28			Street 229, Ahmedabad, Gujarat - 380699	31993			
29			Street 45, Nagpur, Maharashtra - 440765	23250			
30			Street 176, Lucknow, Uttar Pradesh - 226538	74588			
17499			Street 199, Jaipur, Rajasthan - 302715	30183			
	yearly_income	total_debt	credit_score	...	zip	Credit_Tier	\
8258	501564	248128	618	...	400244	Fair	
6665	602667	161562	587	...	560677	Fair	
8747	1424114	382430	586	...	800181	Fair	
8527	505420	99226	611	...	440594	Fair	
2713	787483	429530	597	...	380797	Fair	
...	
27	832468	184888	620	...	411506	Fair	
28	383912	105126	642	...	380699	Fair	
29	278999	0	762	...	440765	Excellent	
30	895055	0	736	...	226538	Good	
17499	362198	0	816	...	302715	Excellent	
	Loan_Default	Fraud	Customer_Segment	Transaction_Type	High_Risk	\	
8258	1	0	2	1	1		
6665	1	0	2	1	1		
8747	1	0	0	0	1		
8527	1	0	2	1	1		
2713	1	0	0	1	1		
...		
27	0	0	3	1	0		
28	0	0	2	0	0		
29	0	0	1	0	0		
30	0	0	3	0	0		
17499	0	0	1	0	0		

	Risk_Band	cum_customers	cum_risk
8258	High Risk	0.00005	0.000348
6665	High Risk	0.00010	0.000695
8747	High Risk	0.00015	0.001043
8527	High Risk	0.00020	0.001390
2713	High Risk	0.00025	0.001738
...
27	Low / Medium Risk	0.99980	1.000000
28	Low / Medium Risk	0.99985	1.000000
29	Low / Medium Risk	0.99990	1.000000
30	Low / Medium Risk	0.99995	1.000000
17499	Low / Medium Risk	1.00000	1.000000

[20000 rows x 30 columns]



Credit risk is highly concentrated, with a relatively small proportion of customers accounting for a disproportionately large share of total portfolio risk.

This concentration indicates that targeted monitoring and stricter controls on a limited high-risk subset can significantly reduce overall portfolio risk, rather than applying uniform policies to all customers.


Q.6::How effectively can customer segmentation be used to design differentiated credit policies or limits?


```
In [50]: # EDA: Average Credit Score & Debt by Customer Segment  
# (to assess differentiated credit policies)
```

```
segment_policy = df.groupby("Customer_Segment")[[  
    "credit_score",  
    "total_debt",  
    "yearly_income"  
]].mean()  
print(segment_policy)  
# Plot 1: Average Credit Score by Segment  
plt.figure(figsize=(8, 5))  
plt.bar(  
    segment_policy.index.astype(str),  
    segment_policy["credit_score"],  
    color=["green", "orange", "blue", "red"]  
)  
  
plt.xlabel("Customer Segment")  
plt.ylabel("Average Credit Score")  
plt.title("Average Credit Score by Customer Segment")  
plt.grid(axis="y", alpha=0.3)  
plt.show()
```

```
# Plot 2: Average Debt by Segment  
plt.figure(figsize=(8, 5))  
plt.bar(  
    segment_policy.index.astype(str),  
    segment_policy["total_debt"],  
    color=["red", "purple", "orange", "green"]  
)  
  
plt.xlabel("Customer Segment")  
plt.ylabel("Average Total Debt")  
plt.title("Average Debt Exposure by Customer Segment")  
plt.grid(axis="y", alpha=0.3)  
plt.show()
```

Customer_Segment	credit_score	total_debt	yearly_income
0	665.243390	479771.795879	1.216282e+06
1	774.487602	59023.090100	4.868866e+05
2	642.353502	99820.300518	4.896814e+05
3	716.288543	88890.426485	9.844204e+05





Customer segmentation clearly enables differentiated credit policies, as Segment 1 shows high credit scores and low debt, making it suitable for higher credit limits and preferential pricing.

Segment 0 and Segment 2 require tighter limits and enhanced monitoring, since high debt levels or lower credit scores indicate elevated risk despite income differences.

Q.7::Which transaction types are most frequently associated with elevated risk levels?

```
In [51]: # Calculate proportion of high-risk customers per transaction type
txn_risk = (
    df.groupby("Transaction_Type")["High_Risk"]
    .mean()
    .sort_values(ascending=False)
)
print(txn_risk)
# Plot
plt.figure(figsize=(8, 5))
txn_risk.plot(
    kind="bar",
    color=["red", "orange", "green"]
)

plt.xlabel("Transaction Type (0 = Low, 1 = Medium, 2 = High)")
```

```
plt.ylabel("Proportion of High-Risk Customers")
plt.title("Credit Risk by Transaction Type")
plt.xticks(rotation=0)
plt.grid(axis="y", alpha=0.3)
plt.show()
```

```
Transaction_Type
2    0.150485
0    0.145767
1    0.142629
Name: High_Risk, dtype: float64
```



High-value transactions (Type 2) show the highest association with elevated credit risk, indicating that larger transaction amounts slightly increase risk exposure.

Risk levels across transaction types are relatively close, suggesting that transaction value alone is not a strong risk driver and should be evaluated alongside customer financial attributes.

Q.8: What is the trade-off between rejecting risky customers and accidentally rejecting creditworthy customers?

```
In [52]: # Confusion matrix using current risk predictions
cm = confusion_matrix(y_test, risk_pred)
```

```

# Extract values
tn, fp, fn, tp = cm.ravel()

tradeoff_df = pd.DataFrame({
    "Decision_Type": [
        "Correctly Approved (TN)",
        "Incorrectly Rejected (FP)",
        "Missed Risky Customer (FN)",
        "Correctly Rejected (TP)"
    ],
    "Count": [tn, fp, fn, tp]
})
print(tradeoff_df)
# Plot trade-off
plt.figure(figsize=(8, 5))
plt.bar(
    tradeoff_df["Decision_Type"],
    tradeoff_df["Count"],
    color=["green", "orange", "red", "darkred"]
)

plt.ylabel("Number of Customers")
plt.title("Trade-off Between Risk Rejection and Customer Rejection")
plt.xticks(rotation=30)
plt.grid(axis="y", alpha=0.3)
plt.show()

```

	Decision_Type	Count
0	Correctly Approved (TN)	5000
1	Incorrectly Rejected (FP)	137
2	Missed Risky Customer (FN)	787
3	Correctly Rejected (TP)	76



The model strongly prioritizes approving creditworthy customers, with a high number of correct approvals and relatively few incorrect rejections.

This conservative strategy results in many missed risky customers, highlighting a clear trade-off between minimizing customer rejection and maximizing risk capture, which can be adjusted through threshold tuning.


Q.9::Can behavioral and financial attributes reliably classify customers into credit score tiers without directly using the credit score?


```
In [54]: tier_behavior = df.groupby("Credit_Tier")[[  
    "yearly_income",  
    "total_debt",  
    "num_credit_cards",  
    "current_age"  
]].mean()  
  
# Plot 1: Average Income by Credit Tier
```


```
plt.figure(figsize=(8, 5))
plt.bar(
    tier_behavior.index,
    tier_behavior["yearly_income"],
    color=["red", "orange", "blue", "green"]
)
plt.xlabel("Credit Score Tier")
plt.ylabel("Average Yearly Income")
plt.title("Income Distribution Across Credit Score Tiers")
plt.grid(axis="y", alpha=0.3)
plt.show()

# Plot 2: Average Debt by Credit Tier
plt.figure(figsize=(8, 5))
plt.bar(
    tier_behavior.index,
    tier_behavior["total_debt"],
    color=["darkred", "gold", "purple", "green"]
)
plt.xlabel("Credit Score Tier")
plt.ylabel("Average Total Debt")
plt.title("Debt Distribution Across Credit Score Tiers")
plt.grid(axis="y", alpha=0.3)
plt.show()

# Plot 3: Credit Cards by Credit Tier
plt.figure(figsize=(8, 5))
plt.bar(
    tier_behavior.index,
    tier_behavior["num_credit_cards"],
    color=["brown", "orange", "blue", "teal"]
)
plt.xlabel("Credit Score Tier")
plt.ylabel("Average Number of Credit Cards")
plt.title("Credit Card Usage Across Credit Score Tiers")
plt.grid(axis="y", alpha=0.3)
plt.show()
```







Q.10::How would changes in risk thresholds affect approval rates, portfolio risk, and business growth?

```
In [55]: risk_probs = risk_prob
y_true = y_test

thresholds = np.arange(0.1, 0.6, 0.05)

approval_rates = []
risk_capture = []

for t in thresholds:
    preds = (risk_probs >= t).astype(int)

    approval_rate = (preds == 0).mean() * 100
    approval_rates.append(approval_rate)

    captured_risk = (preds[y_true == 1] == 1).mean() * 100
    risk_capture.append(captured_risk)
```

```


threshold_df = pd.DataFrame({
    "Threshold": thresholds,
    "Approval_Rate_%": approval_rates,
    "Risk_Capture_%": risk_capture
})
print(threshold_df)

plt.figure(figsize=(8, 5))
plt.plot(threshold_df["Threshold"], threshold_df["Approval_Rate_%"],
          marker="o", color="green", label="Approval Rate")
plt.plot(threshold_df["Threshold"], threshold_df["Risk_Capture_%"],
          marker="o", color="red", label="Risk Capture")

plt.xlabel("Risk Threshold")
plt.ylabel("Percentage (%)")
plt.title("Impact of Risk Thresholds on Approvals & Risk Capture")
plt.legend()
plt.grid(alpha=0.3)
plt.show()

```

	Threshold	Approval_Rate_%	Risk_Capture_%
0	0.10	24.150000	86.326767
1	0.15	63.216667	57.705678
2	0.20	84.966667	28.852839
3	0.25	92.666667	16.106605
4	0.30	96.450000	8.806489
5	0.35	97.516667	7.068366
6	0.40	98.633333	4.055620
7	0.45	98.966667	2.896871
8	0.50	99.400000	1.390498
9	0.55	99.450000	1.158749



Lower risk thresholds significantly increase risk capture but sharply reduce approval rates, indicating aggressive risk control at the cost of business growth.

Higher thresholds maximize approvals but capture very little risk, highlighting a clear trade-off where moderate thresholds (around 0.20–0.30) offer a balanced compromise between growth and risk containment.

Overall Business Recommendations

Adopt a risk-based approval strategy by leveraging risk bands instead of binary approve/reject decisions, enabling higher approval rates while controlling portfolio risk.

Apply differentiated credit policies by customer segment, offering higher limits and preferential terms to low-risk segments while enforcing stricter limits and monitoring for high-risk and over-leveraged segments.

Do not rely on income alone for credit decisions, as high-income customers can still exhibit elevated risk due to high debt exposure; debt-to-income behavior should be prioritized.

Use moderate risk thresholds ($\approx 0.20\text{--}0.30$) to balance business growth and risk control, as overly conservative thresholds significantly reduce approvals while overly lenient thresholds fail to capture meaningful risk.

Implement manual review workflows for medium-risk customers, allowing human judgment to reduce false rejections without materially increasing portfolio risk.

Focus monitoring efforts on a small subset of customers, since risk is highly concentrated and a limited group drives a disproportionate share of potential losses.

Leverage customer segmentation for targeted strategies, such as credit-building programs for emerging customers and enhanced monitoring for high-income, high-debt profiles.

Treat transaction-level classification as a supporting signal, not a standalone risk driver, as transaction value alone shows limited differentiation without customer context.

Use explainable, probability-based risk models as baselines, enabling transparency, regulatory alignment, and easier policy tuning as the business scales.

Conclusion

This project demonstrates how an integrated analytics approach can transform raw customer and transaction data into actionable business decisions. By combining exploratory data analysis, customer segmentation, credit scoring, transaction classification, and probability-based risk modelling, the analysis provides a realistic and explainable view of portfolio risk rather than relying on isolated metrics or high-accuracy claims.

The results show that credit risk is not uniformly distributed across the customer base—risk is concentrated within specific segments and driven more by debt behavior than income alone.

The risk model effectively separates customers into actionable risk bands, enabling balanced approval strategies that support business growth while maintaining control over potential losses.

Customer segmentation further strengthens decision-making by allowing differentiated credit policies instead of one-size-fits-all rules.

Overall, the project highlights the value of explainable, business-aligned analytics in financial decision-making. Rather than optimizing for perfect predictions, the framework prioritizes transparency, risk control, and practical usability, making it well-suited as a baseline system for real-world credit and risk management applications.

In []:

In []:

In []: